

Московский государственный технический университет им. Н.Э. Баумана

Факультет «Информатика и системы управления»

Кафедра «Системы обработки информации и управления»



“Методы машинного обучения”

ЛАБОРАТОРНАЯ РАБОТА № 2. «Изучение библиотек обработки данных»

Студент группы ИУ5-21М

Попков В.Е.

_____ Дата

_____ Подпись

Москва 2019

Часть 1. Выполните первое демонстрационное задание "demo assignment" под названием "Exploratory data analysis with Pandas" со страницы курса <https://mlcourse.ai/assignments>
Объявление библиотек и настройка отображения:

```
In [2]:
import numpy as np
import pandas as pd
pd.set_option('display.max.columns', 100)
import matplotlib.pyplot as plt
import seaborn as sns
```

Подключение предлагаемой выборки и проверка выводом первых 5ти:

```
In [3]:
data = pd.read_csv('C:/jupWork2/data/adult.data.csv' )
data.head()
```

Out[3]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40

Задача 1) Сколько мужчин и женщин отображено в выборке

```
In [4]:
#1. How many men and women (sex feature) are represented in this dataset?
data['sex'].value_counts()
```

Out[4]:

```
Male      21790
Female    10771
```

Name: sex, dtype: int64

Задача 2) Найти средний возраст женщин

```
In [5]:
#2. What is the average age (age feature) of women?
data_fem = data[data['sex']==' Female']
data_fem.head()
data_fem['age'].mean()
```

Out[5]:

```
36.85823043357163
```

Задача 3) Каков процент жителей из Германии?

```
In [6]:
#3. What is the percentage of German citizens (native-country feature)?
```

```
(data['native-country']==' Germany').sum()/data['age'].count()*100
```

Out[6]:

0.42074874850281013

Задачи 4-5) Найти среднее и отклонение возраста для богатых (ЗП>50к) и бедных (ЗП<50к)

In [110]:

```
#4-5. What are the mean and standard deviation of age for those who earn more than 50K per year (salary feature) and those who earn less than 50K per year?
```

```
ages1 = data.loc[data['salary'] == ' >50K', 'age']
```

```
ages2 = data.loc[data['salary'] == ' <=50K', 'age']
```

```
print("rich: {0} +- {1} years, \npoor: {2} +- {3} years.".format(
    round(ages1.mean()), round(ages1.std(), 1),
    round(ages2.mean()), round(ages2.std(), 1)))
```

rich: 44.0 +- 10.5 years,

poor: 37.0 +- 14.0 years.

Задача 6) Проверить, является ли правдой утверждение что те, кто зарабатывает 50к+ имеют образование как минимум выше старшей школы?

In [7]:

```
#6. Is it true that people who earn more than 50K have at least high school education? (education - Bachelors, Prof-school, Assoc-acdm, Assoc-voc, Masters or Doctorate feature)
```

```
#data.loc[data['salary'] == ' >50K', 'education'].unique()
```

```
data.loc[(data['salary'] == ' >50K') &
```

```
(~data['education'].isin([' Bachelors', ' Prof-school', ' Assoc-acdm',
    ' Assoc-voc', ' Masters', ' Doctorate feature']))]
```

Out[7]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	cas los
7	52	Self-emp-not-inc	209642	HS-grad	9	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0
10	37	Private	280464	Some-college	10	Married-civ-spouse	Exec-managerial	Husband	Black	Male	0	0
20	40	Private	193524	Doctorate	16	Married-civ-spouse	Prof-specialty	Husband	White	Male	0	0
27	54	?	180211	Some-college	10	Married-civ-spouse	?	Husband	Asian-Pac-Islander	Male	0	0
38	31	Private	84154	Some-college	10	Married-civ-spouse	Sales	Husband	White	Male	0	0
55	43	Private	237993	Some-college	10	Married-civ-spouse	Tech-support	Husband	White	Male	0	0
63	42	Private	116632	Doctorate	16	Married-civ-spouse	Prof-specialty	Husband	White	Male	0	0
67	53	Private	169846	HS-grad	9	Married-civ-spouse	Adm-clerical	Wife	White	Female	0	0

3612 rows × 15 columns

Задача 7) Отобразить возрастную статистику по каждой расе и полу. Использовать GROUPBY и DESCRIBE. Найти максимальный возраст мужчин Американско-Индийско-Эскимоской расы (?).

In [112]:

```
#7. Display age statistics for each race (race feature) and each gender (sex feature). Use groupby() and describe(). Find the maximum age of men of Amer-Indian-Eskimo race.
```

```
for (race, sex), sub_df in data.groupby(['race', 'sex']):  
    print("Race: {0}, sex: {1}".format(race, sex))  
    print(sub_df['age'].describe())
```

Race: Amer-Indian-Eskimo, sex: Female

```
count    119.000000  
mean      37.117647  
std       13.114991  
min       17.000000  
25%      27.000000  
50%      36.000000  
75%      46.000000  
max       80.000000
```

Name: age, dtype: float64

Race: Amer-Indian-Eskimo, sex: Male

```
count    192.000000  
mean      37.208333  
std       12.049563  
min       17.000000  
25%      28.000000  
50%      35.000000  
75%      45.000000  
max       82.000000
```

Name: age, dtype: float64

Race: Asian-Pac-Islander, sex: Female

```
count    346.000000  
mean      35.089595  
std       12.300845  
min       17.000000  
25%      25.000000  
50%      33.000000  
75%      43.750000  
max       75.000000
```

Name: age, dtype: float64

Race: Asian-Pac-Islander, sex: Male

```
count    693.000000  
mean      39.073593  
std       12.883944  
min       18.000000  
25%      29.000000  
50%      37.000000  
75%      46.000000  
max       90.000000
```

```
Name: age, dtype: float64
Race: Black, sex: Female
count    1555.000000
mean      37.854019
std       12.637197
min       17.000000
25%       28.000000
50%       37.000000
75%       46.000000
max       90.000000
Name: age, dtype: float64
Race: Black, sex: Male
count    1569.000000
mean      37.682600
std       12.882612
min       17.000000
25%       27.000000
50%       36.000000
75%       46.000000
max       90.000000
Name: age, dtype: float64
Race: Other, sex: Female
count    109.000000
mean      31.678899
std       11.631599
min       17.000000
25%       23.000000
50%       29.000000
75%       39.000000
max       74.000000
Name: age, dtype: float64
Race: Other, sex: Male
count    162.000000
mean      34.654321
std       11.355531
min       17.000000
25%       26.000000
50%       32.000000
75%       42.000000
max       77.000000
Name: age, dtype: float64
Race: White, sex: Female
count    8642.000000
mean      36.811618
std       14.329093
min       17.000000
25%       25.000000
50%       35.000000
75%       46.000000
max       90.000000
```

Name: age, dtype: float64

Race: White, sex: Male

count 19174.000000

mean 39.652498

std 13.436029

min 17.000000

25% 29.000000

50% 38.000000

75% 49.000000

max 90.000000

Name: age, dtype: float64

Задача 8) Найти кого больше среди тех кто получает 50к+: женатых или одиноких мужчин

In [113]:

#8. Among whom the proportion of those who earn a lot(>50K) is more: among married or single men (marital-status feature)? Consider married those who have a marital-status starting with Married (Married-civ-spouse, Married-spouse-absent or Married-AF-spouse), the rest are considered bachelors.

```
data.loc[(data['sex'] == ' Male') &
         (data['marital-status'].isin([' Never-married', ' Separated', '
Divorced',
         ' Widowed']))], 'salary'].value_counts()
```

Out[113]:

<=50K 7552

>50K 697

Name: salary, dtype: int64

Такое же сравнение но среди тех кто имеет в статусе "Семейное положение" статус начинающийся с 'married'

In [114]:

```
data.loc[(data['sex'] == ' Male') &
         (data['marital-status'].str.startswith(' Married'))],
'salary'].value_counts()
```

Out[114]:

<=50K 7576

>50K 5965

Name: salary, dtype: int64

Общая статистика по "Семейному положению"

In [115]:

```
data['marital-status'].value_counts()
```

Out[115]:

Married-civ-spouse 14976

Never-married 10683

Divorced 4443

Separated 1025

Widowed 993

Married-spouse-absent 418

Married-AF-spouse 23

Name: marital-status, dtype: int64

Задание 9) Чему равно максимальное число отработанных в неделю часов? Сколько человек так живет и сколько из них получает много?

In [116]:

#9. What is the maximum number of hours a person works per week (hours-per-week feature)? How many people work such a number of hours and what is the percentage of those who earn a lot among them?

```
maxh = data['hours-per-week'].max()
num = data[data['hours-per-week'] == maxh].shape[0]
pers = float(data[(data['hours-per-week'] == maxh)
                  & (data['salary'] == ' >50K')].shape[0]) / num
print("There are {a} men with {b} hpw and there is a percentage of rich among
them {c}%".format(a=num,b=maxh,c=int(100 * pers)))
```

There are 85 men with 99 hpw and there is a percentage of rich among them 29%

In []:

Задание 10) Найти среднее рабочее время относительно получающих больше/меньше 50к по каждой стране

In [117]:

#10. Count the average time of work (hours-per-week) those who earning a little and a lot (salary) for each country (native-country).

```
for (country, salary), sub_df in data.groupby(['native-country', 'salary']):
    print(country, salary, round(sub_df['hours-per-week'].mean(), 2))
```

? <=50K 40.16

? >50K 45.55

Cambodia <=50K 41.42

Cambodia >50K 40.0

Canada <=50K 37.91

Canada >50K 45.64

China <=50K 37.38

China >50K 38.9

Columbia <=50K 38.68

Columbia >50K 50.0

Cuba <=50K 37.99

Cuba >50K 42.44

Dominican-Republic <=50K 42.34

Dominican-Republic >50K 47.0

Ecuador <=50K 38.04

Ecuador >50K 48.75

El-Salvador <=50K 36.03

El-Salvador >50K 45.0

England <=50K 40.48

England >50K 44.53

France <=50K 41.06

France >50K 50.75

Germany <=50K 39.14

Germany >50K 44.98

Greece <=50K 41.81

Greece >50K 50.62

Guatemala <=50K 39.36

Guatemala >50K 36.67

Haiti <=50K 36.33

Haiti >50K 42.75
Holand-Netherlands <=50K 40.0
Honduras <=50K 34.33
Honduras >50K 60.0
Hong <=50K 39.14
Hong >50K 45.0
Hungary <=50K 31.3
Hungary >50K 50.0
India <=50K 38.23
India >50K 46.48
Iran <=50K 41.44
Iran >50K 47.5
Ireland <=50K 40.95
Ireland >50K 48.0
Italy <=50K 39.62
Italy >50K 45.4
Jamaica <=50K 38.24
Jamaica >50K 41.1
Japan <=50K 41.0
Japan >50K 47.96
Laos <=50K 40.38
Laos >50K 40.0
Mexico <=50K 40.0
Mexico >50K 46.58
Nicaragua <=50K 36.09
Nicaragua >50K 37.5
Outlying-US (Guam-USVI-etc) <=50K 41.86
Peru <=50K 35.07
Peru >50K 40.0
Philippines <=50K 38.07
Philippines >50K 43.03
Poland <=50K 38.17
Poland >50K 39.0
Portugal <=50K 41.94
Portugal >50K 41.5
Puerto-Rico <=50K 38.47
Puerto-Rico >50K 39.42
Scotland <=50K 39.44
Scotland >50K 46.67
South <=50K 40.16
South >50K 51.44
Taiwan <=50K 33.77
Taiwan >50K 46.8
Thailand <=50K 42.87
Thailand >50K 58.33
Trinidad&Tobago <=50K 37.06
Trinidad&Tobago >50K 40.0
United-States <=50K 38.8
United-States >50K 45.51
Vietnam <=50K 37.19


```
Vietnam >50K 39.2
Yugoslavia <=50K 41.6
Yugoslavia >50K 49.5
```

Часть 2. Выполните следующие запросы с использованием двух различных библиотек - Pandas и PandaSQL:

-один произвольный запрос на соединение двух наборов данных -один произвольный запрос на группировку набора данных с использованием функций агрегирования Сравните время выполнения каждого запроса в Pandas и PandaSQL.

Подключаем библиотеки:

```
In [74]:
import numpy as np
import pandas as pd
import pandasql as ps
from timeit import default_timer as timer
```

Датафреймы для двух выборок:

```
In [75]:
fd = pd.read_csv('C:/jupWork2/data/user_device.csv')
sd = pd.read_csv('C:/jupWork2/data/user_usage.csv')
```

Часть выборки User_device:

```
In [76]:
fd.head()
```

Out[76]:

	use_id	user_id	platform	platform_version	device	use_type_id
0	22782	26980	ios	10.2	iPhone7,2	2
1	22783	29628	android	6.0	Nexus 5	3
2	22784	28473	android	5.1	SM-G903F	1
3	22785	15200	ios	10.2	iPhone7,2	3
4	22786	28239	android	6.0	ONE E1003	1

Часть выборки User_usege:

```
In [77]:
sd.head()
```

Out[77]:

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id
0	21.97	4.82	1557.33	22787
1	1710.08	136.88	7267.55	22788
2	1710.08	136.88	7267.55	22789
3	94.46	35.17	519.12	22790
4	71.59	79.26	1557.33	22792

Как видно выше, имеется общее поле use_id. Реализуем соединение двух таблиц по этому полю с использованием PandaSQL:

In [78]:

```
def example1_pandasql(fd, sd):  
  
    simple_query = '''  
        SELECT *  
        FROM fd JOIN sd  
        WHERE fd.use_id==sd.use_id  
        '''  
  
    return ps.sqldf(simple_query, locals())
```

In [79]:

```
example1_pandasql(fd, sd)
```

Out[79]:

	use_id	user_id	platform	platform_version	device	use_type_id	outgoing_mins_per_month	outgoing_sms
0	22787	12921	android	4.3	GT-I9505	1	21.97	4.82
1	22788	28714	android	6.0	SM-G930F	1	1710.08	136.88
2	22789	28714	android	6.0	SM-G930F	1	1710.08	136.88
3	22790	29592	android	5.1	D2303	1	94.46	35.17
4	22792	28217	android	5.1	SM-G361F	1	71.59	79.26
5	22793	28217	android	5.1	SM-G361F	1	71.59	79.26
6	22794	28217	android	5.1	SM-G361F	1	71.59	79.26
7	22795	28217	android	5.1	SM-G361F	1	71.59	79.26
8	22799	29643	android	6.0	ONEPLUS A3003	1	30.92	22.77
9	22801	10976	android	4.4	GT-I9505	1	69.80	14.70
10	22804	29645	android	6.0	SM-G935F	1	554.41	150.06
11	22805	29646	android	4.2	GT-I9195	1	189.10	24.08
12	22806	21615	android	6.0	A0001	1	283.30	107.47
13	22808	29065	android	6.0	SM-G900F	1	324.34	92.52
14	22813	23415	android	4.4	HTC Desire 510	1	797.06	7.67
15	22814	23415	android	4.4	HTC Desire 510	1	797.06	7.67
16	22815	23415	android	4.4	HTC Desire 510	1	797.06	7.67

159 rows × 10 columns

И Pandas. Замерим время выполнения каждого:

In [80]:

```
t1 = timer()  
example1_pandasql(fd, sd)  
elapsed1 = timer() - t1  
t2 = timer()  
fsd = fd.merge(sd)  
elapsed2 = timer() - t2  
print("PandasSQL: {a} \nPandas: {b} ".format(a=elapsed1,b=elapsed2))
```

PandasSQL: 0.042729509363198304

Pandas: 0.01557342086925928

Для полученной соединенной таблице попробуем провести агрегирование с группировкой и так же замерим время выполнения: PandaSQL:

In [81]:

```
def example2_pandasql(sd):  
    aggr_query = '''  
        SELECT distinct  
            device, avg(monthly_mb) as avg_mb  
        FROM sd  
        GROUP BY device  
    '''  
    return ps.sqldf(aggr_query, locals())
```

In [82]:

```
example2_pandasql(fsd)
```

Out[82]:

	device	avg_mb
0	A0001	15573.330000
1	C6603	1557.330000
2	D2303	519.120000
3	D5503	1557.330000
4	D5803	1557.330000
5	D6603	7267.550000
6	E6653	5191.120000
7	EVA-L09	1557.330000
8	F3111	2076.450000
9	GT-I8190N	407.010000
10	GT-I9195	1211.260000
11	GT-I9300	464.185000
12	GT-I9505	5564.726364
13	GT-I9506	803.240000
14	GT-I9515	1557.330000
15	GT-N7100	11939.560000
16	HTC Desire 510	12562.488000
17	HTC Desire 530	1557.330000
18	HTC Desire 620	74.400000
19	HTC Desire 626	519.120000
20	HTC Desire 825	5498.970000

20	HTC Desire 825	5498.970000
21	HTC One M9	2362.070000
22	HTC One S	1038.210000
23	HTC One mini 2	13842.956667
24	HTC One_M8	6577.120000
25	HUAWEI CUN-L01	11.680000
26	HUAWEI VNS-L31	3114.670000
27	LG-H815	1557.330000
28	Lenovo K51c78	1557.330000
29	Moto G (4)	5191.120000
30	MotoE2(4G-LTE)	212.640000
31	Nexus 5X	1557.330000
32	ONE A2003	2076.450000
33	ONEPLUS A3003	3823.610000
34	SM-A300FU	1687.112500
35	SM-A310F	1557.330000
36	SM-A500FU	1557.330000
37	SM-G360F	1557.330000
38	SM-G361F	934.404000
39	SM-G531F	2076.450000
40	SM-G800F	1557.330000

41	SM-G900F	3841.427333
42	SM-G903F	1557.330000
43	SM-G920F	1985.168000
44	SM-G925F	3633.775000
45	SM-G930F	7959.700000
46	SM-G935F	4568.182000
47	SM-J320FN	830.574000
48	SM-N9005	16611.550000
49	SM-N910F	8038.370000
50	VF-795	1557.330000
51	Vodafone Smart ultra 6	5191.120000
52	X11	12458.670000
53	iPhone6,2	650.920000
54	iPhone7,2	1271.390000

Аналогичная функция на Pandas и подсчет времени выполнения.

In [83]:

```
t1 = timer()
example2_pandasql(fsd)
elapsed1 = timer() - t1
t2 = timer()
fsd.groupby('device').monthly_mb.mean()
elapsed2 = timer() - t2
print("PandasSQL: {a} \nPandas: {b} ".format(a=elapsed1,b=elapsed2))

PandasSQL: 0.025251644065065193
Pandas: 0.004341345082139014
```

Как можно видеть в обоих случаях Pandas превосходит PandaSQL почти на порядок. Несмотря на то что цифры в данном примере выглядят незначительными, для огромных датасетов параметр времени может стать значительным. Так же видно еще одно преимущество Pandas - размеры кода. То что реализовано на PandaSQL полноценной функцией-запросом пишется на Pandas в одну строчку. Таким образом за исключением совсем сложных SQL-запросов Pandas будет и быстрее, и короче.